

Stochastic Gradient Method

V. Leclère (ENPC)

June 7th, 2024

Why should I bother to learn this stuff?

- Main algorithm principle for training machine learning model, and in particular deep neural network
- \implies useful for
 - ▶ understanding how the library train ML models
 - ▶ specialization in optimization
 - ▶ specialization in machine learning

Contents

- 1 Setting up the problem
- 2 Full batch method
- 3 Stochastic and minibatch version
- 4 Wrap-up

The optimization problem



We consider the following optimization problem

$$\underset{x \in \mathbb{R}^p}{\text{Min}} \quad F(x) := \mathbb{E} \left[f(x, \xi) \right]$$

where ξ is a random variable.

Contents

- 1 Setting up the problem
- 2 Full batch method**
- 3 Stochastic and minibatch version
- 4 Wrap-up



$$F(\mathbf{x}) := \mathbb{E} \left[f(\mathbf{x}, \boldsymbol{\xi}) \right]$$

Under some regularity conditions (e.g. $f(\cdot, \xi)$ differentiable, $\frac{\partial f(\mathbf{x}, \cdot)}{\partial \mathbf{x}}$ Lipschitz, and $\boldsymbol{\xi}$ integrable) we have

$$\nabla F(\mathbf{x}) = \mathbb{E} \left[\frac{\partial f}{\partial \mathbf{x}}(\mathbf{x}, \boldsymbol{\xi}) \right]$$

This is obvious if $\boldsymbol{\xi}$ is finitely supported : $\text{supp}(\boldsymbol{\xi}) = \{\xi_i\}_{i \in [M]}$, and $p_i := \mathbb{P}(\boldsymbol{\xi} = \xi_i)$,

$$\nabla F(\mathbf{x}) = \frac{\partial}{\partial \mathbf{x}} \left(\sum_{i \in [M]} p_i f(\mathbf{x}, \xi_i) \right) = \sum_{i \in [M]} p_i \frac{\partial}{\partial \mathbf{x}} f(\mathbf{x}, \xi_i)$$



$$F(\mathbf{x}) := \mathbb{E} \left[f(\mathbf{x}, \boldsymbol{\xi}) \right]$$

Under some regularity conditions (e.g. $f(\cdot, \xi)$ differentiable, $\frac{\partial f(\mathbf{x}, \cdot)}{\partial \mathbf{x}}$ Lipschitz, and $\boldsymbol{\xi}$ integrable) we have

$$\nabla F(\mathbf{x}) = \mathbb{E} \left[\frac{\partial f}{\partial \mathbf{x}}(\mathbf{x}, \boldsymbol{\xi}) \right]$$

This is obvious if $\boldsymbol{\xi}$ is finitely supported : $\text{supp}(\boldsymbol{\xi}) = \{\xi_i\}_{i \in [M]}$, and $p_i := \mathbb{P}(\boldsymbol{\xi} = \xi_i)$,

$$\nabla F(\mathbf{x}) = \frac{\partial}{\partial \mathbf{x}} \left(\sum_{i \in [M]} p_i f(\mathbf{x}, \xi_i) \right) = \sum_{i \in [M]} p_i \frac{\partial}{\partial \mathbf{x}} f(\mathbf{x}, \xi_i)$$

Standard continuous optimization method

Thus, we are looking at

$$\min_{x \in \mathbb{R}^p} F(x)$$

where F is a (strongly) convex differentiable function if $f(\cdot, \xi)$ is, and we know how to compute its gradient.

Thus, we should be able to solve our problem through the method presented in earlier courses:

- gradient algorithm
- conjugate gradient
- Newton / Quasi-Newton

Why bother with another (class of) algorithm ?

Standard continuous optimization method

Thus, we are looking at

$$\min_{x \in \mathbb{R}^p} F(x)$$

where F is a (strongly) convex differentiable function if $f(\cdot, \xi)$ is, and we know how to compute its gradient.

Thus, we should be able to solve our problem through the method presented in earlier courses:

- gradient algorithm
- conjugate gradient
- Newton / Quasi-Newton

Why bother with another (class of) algorithm ?

Standard continuous optimization method

Thus, we are looking at

$$\min_{x \in \mathbb{R}^p} F(x)$$

where F is a (strongly) convex differentiable function if $f(\cdot, \xi)$ is, and we know how to compute its gradient.

Thus, we should be able to solve our problem through the method presented in earlier courses:

- gradient algorithm
- conjugate gradient
- Newton / Quasi-Newton

Why bother with another (class of) algorithm ?

Contents

- 1 Setting up the problem
- 2 Full batch method
- 3 Stochastic and minibatch version
- 4 Wrap-up

Computing the gradient is costly



For a given solution $\mathbf{x} \in \mathbb{R}^p$ computing the gradient

$$\nabla F(\mathbf{x}) = \mathbb{E} \left[\frac{\partial f(\mathbf{x}, \boldsymbol{\xi})}{\partial \mathbf{x}} \right]$$

is costly as it requires computing a multidimensional integral (if $\boldsymbol{\xi}$ admits a density), or a large sum.

Indeed, in most machine learning applications, we consider that $\boldsymbol{\xi}$ is uniformly distributed over the data (*empirical risk minimization*), thus computing the gradient requires a pass over every sample in the dataset.

Datasets of size $N > 10^6$ are common.

Computing the gradient is costly



For a given solution $\mathbf{x} \in \mathbb{R}^p$ computing the gradient

$$\nabla F(\mathbf{x}) = \mathbb{E} \left[\frac{\partial f(\mathbf{x}, \boldsymbol{\xi})}{\partial \mathbf{x}} \right]$$

is costly as it requires computing a multidimensional integral (if $\boldsymbol{\xi}$ admits a density), or a large sum.

Indeed, in most machine learning applications, we consider that $\boldsymbol{\xi}$ is uniformly distributed over the data (*empirical risk minimization*), thus computing the gradient requires a pass over every sample in the dataset.

Datasets of size $N > 10^6$ are common.

Computing the gradient is costly



For a given solution $\mathbf{x} \in \mathbb{R}^p$ computing the gradient

$$\nabla F(\mathbf{x}) = \mathbb{E} \left[\frac{\partial f(\mathbf{x}, \boldsymbol{\xi})}{\partial \mathbf{x}} \right]$$

is costly as it requires computing a multidimensional integral (if $\boldsymbol{\xi}$ admits a density), or a large sum.

Indeed, in most machine learning applications, we consider that $\boldsymbol{\xi}$ is uniformly distributed over the data (*empirical risk minimization*), thus computing the gradient requires a pass over every sample in the dataset.

Datasets of size $N > 10^6$ are common.

Estimating the gradient



Instead of using a true gradient

$$g^{(k)} = \nabla F(x^{(k)})$$

we can use a *statistical estimator* of the gradient

$$\hat{g}^{(k)} \rightsquigarrow g^{(k)} = \mathbb{E} \left[\frac{\partial f(x^{(k)}, \xi)}{\partial x} \right]$$

The most standard estimator being

$$\hat{g}^{(k)} = \frac{\partial f(x^{(k)}, \xi^{(k)})}{\partial x}$$

where $\xi^{(k)}$ is drawn randomly according to the law of ξ (i.e. it is a random datapoint).

Estimating the gradient



Instead of using a true gradient

$$g^{(k)} = \nabla F(x^{(k)})$$

we can use a *statistical estimator* of the gradient

$$\hat{g}^{(k)} \rightsquigarrow g^{(k)} = \mathbb{E} \left[\frac{\partial f(x^{(k)}, \xi)}{\partial x} \right]$$

The most standard estimator being

$$\hat{g}^{(k)} = \frac{\partial f(x^{(k)}, \xi^{(k)})}{\partial x}$$

where $\xi^{(k)}$ is drawn randomly according to the law of ξ (i.e. it is a random datapoint).



Pros:

- computing $\hat{g}^{(k)} = \frac{\partial f(x^{(k)}, \xi^{(k)})}{\partial x}$ is really easy
- we do not need to spend lots of time early to get a precise gradient
- we can stop whenever we want (do not need a full pass on the data)

Cons:

- $\hat{g}^{(k)}$ is a noisy estimator of the gradient
- requires a new convergence theory
- $x^{(k+1)} := x^{(k)} - \alpha \hat{g}^{(k)}$ generally does not converge almost surely to the optimal solution as this makes a noisy trajectory



- At optimality we should have $\nabla F(x^\sharp) = 0$
- It doesnot mean that $\frac{\partial f(x^\sharp, \xi^{(k)})}{\partial x}$ equals 0 !
- In particular there is no reason for $\hat{g}^{(k)}$ to be eventually small, only its expectation should be small!

• \leadsto we generally use either:

- ▶ decreasing step e.g. $\alpha^{(k)} = \frac{\alpha^{(0)}}{k}$
- ▶ average points $\bar{x}^{(k)} = \frac{1}{k} \sum_{\kappa \leq k} x^{(\kappa)}$



- At optimality we should have $\nabla F(x^\sharp) = 0$
- It doesnot mean that $\frac{\partial f(x^\sharp, \xi^{(k)})}{\partial x}$ equals 0 !
- In particular there is no reason for $\hat{g}^{(k)}$ to be eventually small, only its expectation should be small!
- \leadsto we generally use either:
 - ▶ decreasing step e.g. $\alpha^{(k)} = \frac{\alpha^{(0)}}{k}$
 - ▶ average points $\bar{x}^{(k)} = \frac{1}{k} \sum_{\kappa \leq k} x^{(\kappa)}$



- $\hat{g}^{(k)} = \frac{\partial f(x^{(k)}, \xi^{(k)})}{\partial x}$ is an easy-to-compute but noisy estimator of the gradient
- $\hat{g}^{(k)} = \frac{1}{N} \sum_{i \in [N]} \frac{\partial f(x^{(k)}, \xi_i)}{\partial x}$ is a long (full batch) to compute but perfect estimator
- minibatch aims at a middle ground: randomly draw a sample S of realizations of ξ , and use $\hat{g}^{(k)} = \frac{1}{|S|} \sum_{\xi \in S} \frac{\partial f(x^{(k)}, \xi)}{\partial x}$

Video explanation

A video explaining why the stochastic gradient might be efficient

<https://www.youtube.com/watch?v=UmathvAKj80> (3')

Some adaptations of the stochastic gradient method are presented in the following video: <https://www.youtube.com/watch?v=NE88eqLngkg> (15')

Contents

- 1 Setting up the problem
- 2 Full batch method
- 3 Stochastic and minibatch version
- 4 Wrap-up**

What you have to know

- That for a stochastic problem gradient step requires to compute an expectation
- That stochastic gradient do not compute the true gradient, but only an estimator of the gradient

What you really should know

- gradient algorithm (or more advanced version) is faster in term of number of iterations
- stochastic gradient needs more iteration, but each is faster

What you have to be able to do

- dive in the scientific litterature on the subject if you need to implement this type of algorithm